

# Week 6: Mixture Models

PLSC 40502 - Statistical Models

Review

# Previously

- **Survey weighting**
  - Learning about **population** parameters from non-representative samples.
  - Two key components: **selection/non-response model** and a **measurement model** for the population targets.
  - **Calibration** weighting - find weights that satisfy a set of moment condition subject to a loss function.
- **Multilevel Regression and Post-stratification**
  - Can we estimate population proportions for small areas with a national survey?
  - **Yes** - if we pool information from other units.

# This week

- **"Exploratory" modeling**
  - Unsupervised "clustering" algorithms
  - K-means clustering
  - Bayesian mixture models
  - Obtaining MLEs or posterior modes via Expectation-Maximization
- **Topic modeling**
  - A mixture model for text:
  - Documents have a topic **distribution**
  - Define a DGP for each word as a function of the document topic distribution and a topic-specific word distribution.
  - Flexible way of representing documents via a lower-dimensional summary.

# Finite Mixture Models

# Finite Mixture Models

- In conventional **regression** models, we assume a known distribution for the outcome given a set of covariates
  - The covariates could be something like a discrete "group" label (e.g.  $X_i = x$ )
- In this setting, conditional on  $X_i$ , we would model each  $Y_i$  with a known distribution (e.g. Normal)

$$Y_i | X_i = x \sim \text{Normal}(\mu_x, \sigma_x^2)$$

$\mu_x$  and  $\sigma_x^2$  are the mean and variance parameters associated with the group  $X_i = x$ .

- Now imagine that the group indicators are **not observed** but instead rather **latent parameters**
  - We can do inference on these parameters conditional on the observed data and the model.

# Finite Mixture Models

- In a **finite** mixture model, we assume that the outcome distribution for each observation  $i$  is governed by some discrete latent indicator  $z_i \in \{1, 2, \dots, K\}$ 
  - $z_1, z_2, \dots, z_N$  are i.i.d.  $\sim \text{Categorical}(\pi)$
- Conditional on the latent variable  $z_i = k$ , the outcome distribution is known.
- For example, in a **gaussian** finite mixture model, we assume

$$Y_i | z_i = k \sim \text{Normal}(\mu_k, \sigma_k^2)$$

- But unconditionally,  $Y_i$  has a **mixture** distribution in that its density is a weighted average of the component densities.
- In a **gaussian** finite mixture model, we have:

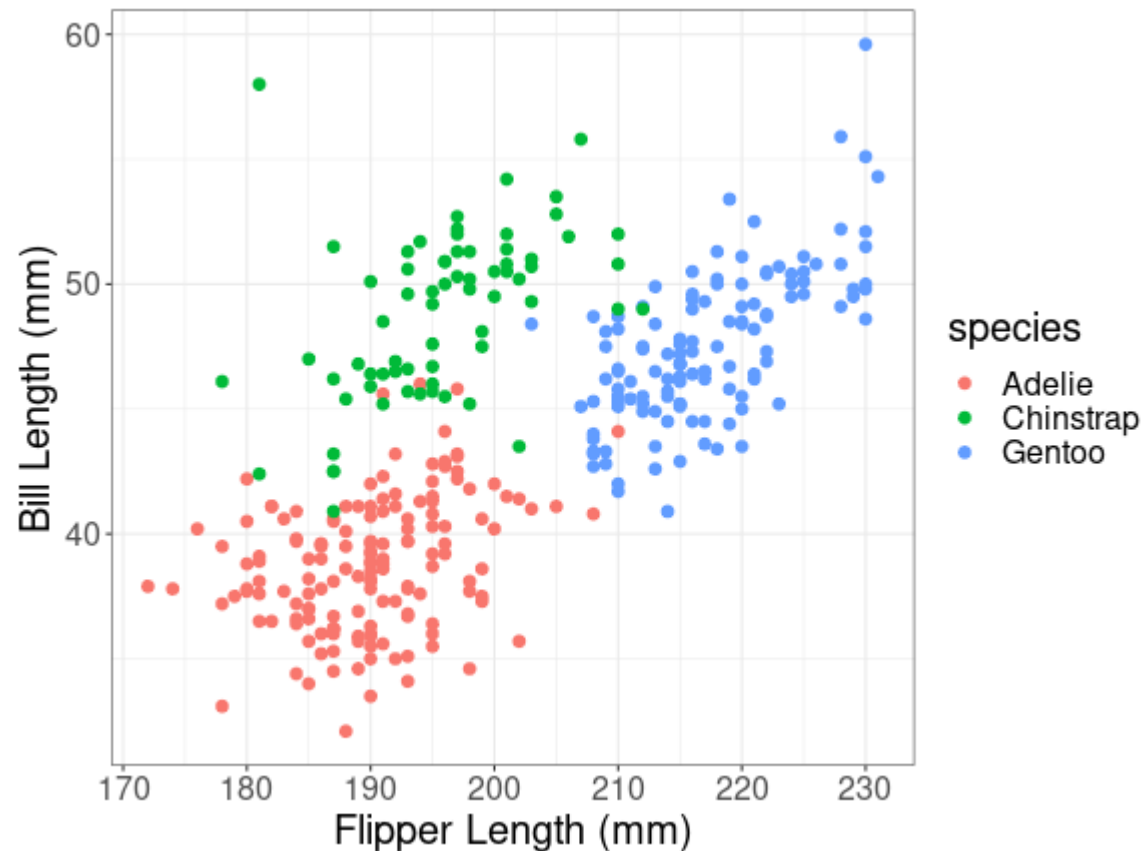
$$Y_i \sim \sum_{k=1}^K \pi_k \times \text{Normal}(\mu_k, \sigma_k^2)$$

where  $\pi_k = \text{Pr}(z_i = k)$

- Our goal is to estimate the cluster means/variances for each cluster and the mixing proportions using likelihood inference.
  - Put priors on the mean/variance parameters and  $\pi$  to make it fully bayesian

# Penguins!


- For our running example, we'll look at the `palmerpenguins` dataset which contains measurements on three species of penguins in the Palmer Archipelago in Antarctica





# Clustering

- Suppose we didn't observe the labels, could we still recover the latent "clusters" in the data?




# K-means

- A simple algorithm for generating clusters that has no underlying probabilistic model is the **K-means algorithm**
  - Straightforward to implement and still surprisingly popular and effective for a "first cut" at the data.
- **K-means algorithm**
  1. Initialize  $K$  distinct clusters by randomly assigning points to one of the  $k$  groups
  2. Calculate  $\mu_k$  as the mean of  $Y_i$  in observations in cluster  $k$ .
  3. Reassign each point to the cluster  $k$  that has the smallest distance between  $Y_i$  and  $\mu_k$ .
  4. Repeat 2-3 until no points change their assignments.


# K-means example

- Let  $K = 3$  and randomly pick three initial  $\mu_k$




# K-means example

- Assign each penguin to the cluster with the nearest cluster centroid.




# K-means example

- Recompute the cluster means




# K-means example

- Re-label the points




# K-means example

- Iterate until convergence



# K-means example

- Note some sensitivity to the choice of distance metric. Here's what it looks like if we standardize each  $X$  by its standard deviation.





# K-means example

- How do our clusters compare to the **actual** categories? Surprisingly well!

```
table(penguins_complete$cluster2, penguins_complete$species)
```

```
##  
##      Adelie Chinstrap Gentoo  
##  1      146          5       0  
##  2         1          4      122  
##  3         4         59       1
```

- But the clusters themselves don't have any underlying interpretation or meaning
  - We **assign** the interpretation to them through inspection!

# K-means

- Advantages
  - Easy to implement, fast updating steps and quick convergence.
- Disadvantages
  - Multimodality (we can't really solve this one!)
  - Choice of distance metric matters.
  - Hard to make sense of outcomes that aren't in  $\mathbb{R}^d$  (e.g. binary/discrete outcomes)
- With conventional K-means, we have no underlying probability model -- each unit is assigned to a single cluster.
  - Sometimes called **"hard"** K-means.
- We can instead implement an approach sometimes referred to as a "soft" K-means algorithm
  - Our target of inference is the **probability** that a unit belongs to each cluster.

# Gaussian Mixture Models

- Assume **latent variables**  $z = \{z_1, z_2, \dots, z_N\}$

$$z_i \sim \text{i.i.d. Categorical}(\pi)$$

- Then, the outcome vector  $Y_i$  has a multivariate gaussian distribution conditional on  $Z_i = k$

$$Y_i | z_i = k \sim \text{Normal}(\mu_k, \Sigma_k)$$

- Our goal is to infer  $\mu_k$  and  $\Sigma_k$  for each cluster
  - And conditional on these estimates, we can obtain the cluster probabilities  $Pr(Z_i = k | Y_i)$

# Expectation-Maximization

- Let's write down the marginal likelihood, marginalizing out the latent parameters:

$$\mathcal{L}(\mu, \Sigma, \pi; \mathbf{Y}) = \prod_{i=1}^N \sum_{k=1}^K f(Y_i | \mu_k, \Sigma_k, Z_i = k) \times p(Z_i = k | \pi)$$

- This is a tricky likelihood to maximize - if we take the log, we get a log of a sum (which doesn't simplify as neatly as a log of a product)

$$\ell(\mu, \Sigma, \pi; \mathbf{Y}) = \sum_{i=1}^N \log \left( \sum_{k=1}^K \pi_k \times \mathcal{N}(Y_i | \mu_k, \Sigma_k) \right)$$

- Suppose we knew the  $Z_i$  as well (they were like "data"), then the "complete" log-likelihood would be easier to maximize!

$$\ell(\mu, \Sigma, \pi; \mathbf{Y}, \mathbf{Z}) = \sum_{i=1}^N \sum_{k=1}^K \log \left( \pi_k^{I(Z_i=k)} \times \mathcal{N}(Y_i | \mu_k, \Sigma_k)^{I(Z_i=k)} \right)$$

# Expectation-Maximization

- The **Expectation-Maximization** algorithm is another iterative method for obtaining a maximum likelihood estimate (or maximum a posteriori (MAP) in a Bayesian setting) when the likelihood consists of a sum/integral over the distribution of some **latent variables**
  - It uses the idea of iteratively optimizing a **lower bound** on the likelihood until convergence is reached.
  - The key trick is **Jensen's inequality**. For a **concave** function like the logarithm:

$$\log(E[X]) \geq E[\log(X)]$$

- Consider the general setting where we have a parameter  $\theta$ , data  $\mathbf{X}$  and discrete latent variables  $Z$
- We want to optimize the log-likelihood:

$$\ell(\theta; \mathbf{X}) = \log \sum_{\mathbf{Z}} f(\mathbf{X}, \mathbf{Z} | \theta) = \log \sum_{\mathbf{Z}} f(\mathbf{X} | \mathbf{Z}, \theta) p(\mathbf{Z} | \theta)$$

# Expectation-Maximization

- Instead of optimizing the log-likelihood, let's try to come up with a lower-bound. Define an arbitrary distribution on the latent variables  $q(\mathbf{Z})$ .
- Multiply by 1:

$$\ell(\theta; \mathbf{X}) = \log \sum_{\mathbf{Z}} f(\mathbf{X} | \mathbf{Z}, \theta) p(\mathbf{Z} | \theta) \frac{q(\mathbf{Z})}{q(\mathbf{Z})}$$

- Rearranging terms, we can see that this can be written as an expectation over the distribution  $q$

$$\ell(\theta; \mathbf{X}) = \log E_q \left[ \frac{f(\mathbf{X} | \mathbf{Z}, \theta) p(\mathbf{Z} | \theta)}{q(\mathbf{Z})} \right]$$

- By **Jensen's inequality**, we have a lower bound

$$\ell(\theta; \mathbf{X}) \geq E_q \left[ \log \left( \frac{f(\mathbf{X} | \mathbf{Z}, \theta) p(\mathbf{Z} | \theta)}{q(\mathbf{Z})} \right) \right]$$

# Expectation-Maximization

- And by properties of logs, our lower-bound is:

$$\ell(\theta; \mathbf{X}) \geq E_q[\log(f(\mathbf{X} | \mathbf{Z}, \theta))] + E_q[\log(f(\mathbf{Z} | \theta))] - E_q[\log(q(\mathbf{Z}))]$$

- We can optimize this **iteratively** by switching between finding an optimal distribution  $q^{(t+1)}$  given parameter values  $\theta^{(t)}$  and finding parameter values  $\theta^{(t+1)}$  given an existing choice of  $q^{(t)}$ .
  1. **Expectation** step: Find  $q^{(t+1)}$  (what distribution are we taking the expectation over)
  2. **Maximization** step: Find  $\theta^{(t+1)}$  (given our  $q$  distribution, what is the value of the parameter values that maximizes the lower bound).

# E-step

- Can we find a closed-form "optimal" solution for our update  $q^{(t+1)}$  given  $\theta^{(t)}$ ?
  - Yes, find the  $q$  that makes the inequality an equality!

$$\ell(\theta^{(t)}; \mathbf{X}) \geq E_q \left[ \log \left( \frac{f(\mathbf{X} | \mathbf{Z}, \theta^{(t)}) p(\mathbf{Z} | \theta^{(t)})}{q(\mathbf{Z})} \right) \right]$$

$$\ell(\theta^{(t)}; \mathbf{X}) \geq \sum_{\mathbf{Z}} \log \left( \frac{f(\mathbf{X} | \mathbf{Z}, \theta^{(t)}) p(\mathbf{Z} | \theta^{(t)})}{q(\mathbf{Z})} \right) q(\mathbf{Z})$$



# E-step

- Let's try to get the  $\mathbf{Z}$  out of the log. Start by factoring the joint distribution in the numerator conditioning on  $\mathbf{X}$  instead of  $\mathbf{Z}$

$$\ell(\theta^{(t)}; \mathbf{X}) \geq \sum_{\mathbf{Z}} \log \left( \frac{f(\mathbf{Z} | \mathbf{X}, \theta^{(t)}) p(\mathbf{X} | \theta^{(t)})}{q(\mathbf{Z})} \right) q(\mathbf{Z})$$

- Now we can see the optimal choice for  $q^{t+1}$  revealed to us. Suppose we plug in  $q = f(\mathbf{Z} | \mathbf{X}, \theta^{(t)})$ , the conditional distribution of  $\mathbf{Z}$  given  $\mathbf{X}$  and  $\theta^{(t)}$ :

$$\ell(\theta^{(t)}; \mathbf{X}) \geq \sum_{\mathbf{Z}} \log(p(\mathbf{X} | \theta^{(t)}) \times f(\mathbf{Z} | \mathbf{X}, \theta^{(t)}))$$

- The logged term no longer depends on  $\mathbf{Z}$ , so pull it out of the sum

$$\ell(\theta^{(t)}; \mathbf{X}) \geq \log(p(\mathbf{X} | \theta^{(t)})) \times \sum_{\mathbf{Z}} f(\mathbf{Z} | \mathbf{X}, \theta^{(t)})$$

# E-step

- The sum is equal to 1 (sum over a PMF/PDF) and we're left with the definition of the marginal log-likelihood, so this expression holds with *equality*

$$\ell(\theta^{(t)}; \mathbf{X}) = \log(p(\mathbf{X} | \theta^{(t)}))$$

- So our choice of distribution over which to take the expectation of the joint likelihood is  $q^{(t+1)} = f(\mathbf{Z} | \mathbf{X}, \theta^{(t)})$

# M-step

- Let's go back to our lower bound - for a given value of  $q^{(t)}$ , we want to find the  $\theta^{(t+1)}$  that maximize the "complete data" log-likelihood

$$\ell(\theta; \mathbf{X}) \geq E_{q^{(t)}}[\log(f(\mathbf{X} | \mathbf{Z}, \theta))] + E_{q^{(t)}}[\log(f(\mathbf{Z} | \theta))] - E_{q^{(t)}}[\log(q^{(t)}(\mathbf{Z}))]$$

- Since the third term doesn't depend on  $\theta$ , we only need to worry about the first two.
  - This is sometimes called the "Q-Function"  $Q(\theta | \theta^{(t)})$  or the expected "complete data" log-likelihood

$$Q(\theta | \theta^{(t)}) = E_{\mathbf{Z} | \mathbf{X}, \theta^{(t)}}[\log(f(\mathbf{X}, \mathbf{Z} | \theta))]$$

- The M-step sets  $\theta^{(t+1)}$  to the value of  $\theta$  that maximizes this Q-function

$$\theta^{(t+1)} = \arg \max_{\theta} Q(\theta | \theta^{(t)})$$

# Deriving EM for the GMM

- For current values of  $\mu^{(t)}$ ,  $\Sigma^{(t)}$ ,  $\pi^{(t)}$ , let's derive the  $Q$  function.
- We'll start by deriving the conditional distribution  $z_i | Y_i, \mu^{(t)}, \Sigma^{(t)}, \pi^{(t)}$  using Bayes' rule

$$p(z_i = k | Y_i, \mu^{(t)}, \Sigma^{(t)}, \pi^{(t)}) = \frac{p(z_i = k | \pi_k^{(t)}) f(Y_i | \mu_k^{(t)}, \Sigma_k^{(t)}, z_i = k)}{f(Y_i | \mu^{(t)}, \Sigma^{(t)})}$$

- Given  $z_i = k$ , we know the distribution is normal at mean  $\mu_k^{(t)}$  and variance  $\Sigma_k^{(t)}$ . And the denominator is just the marginal

$$\gamma_i^k = p(z_i = k | Y_i, \mu^{(t)}, \Sigma^{(t)}, \pi^{(t)}) = \frac{\pi_k^{(t)} \times \mathcal{N}(Y_i | \mu_k^{(t)}, \Sigma_k^{(t)})}{\sum_{j=1}^K \pi_j^{(t)} \times \mathcal{N}(Y_i | \mu_j^{(t)}, \Sigma_j^{(t)})}$$

- These weights,  $\gamma_i^k$  are sometimes called the "responsibility" parameters as they denote the extent to which each cluster is "responsible" for an observation.

# Deriving EM for the GMM

- Recall the "complete likelihood"

$$\ell(\mu, \Sigma, \pi; \mathbf{Y}, \mathbf{Z}) = \sum_{i=1}^N \sum_{k=1}^K \log \left( \pi_k^{I(Z_i=k)} \times \mathcal{N}(Y_i | \mu_k, \Sigma_k)^{I(Z_i=k)} \right)$$

- Simplify it a bit

$$\ell(\mu, \Sigma, \pi; \mathbf{Y}, \mathbf{Z}) = \sum_{i=1}^N \sum_{k=1}^K I(Z_i = k) \log(\pi_k) + \sum_{i=1}^N \sum_{k=1}^K I(Z_i = k) \log \mathcal{N}(Y_i | \mu_k, \Sigma_k)$$

- Now, taking the expectation over  $Z_i$ , the only component that is not a constant is  $I(Z_i = k)$  and  $E[I(Z_i = k)] = p(Z_i = k)$  (fundamental bridge).
  - And we got that (conditional) expectation in the previous section:  $\gamma_i^k$
- So our Q function is

$$Q(\theta | \theta^{(t)}) = \sum_{i=1}^N \sum_{k=1}^K \gamma_i^k \log(\pi_k) + \sum_{i=1}^N \sum_{k=1}^K \gamma_i^k \log \mathcal{N}(Y_i | \mu_k, \Sigma_k)$$

# Deriving EM for the GMM

- Closed form solutions are straightforward to obtain for the M-step (and follow from weighted regression)

$$n_k = \sum_{i=1}^N \gamma_i^k$$

$$\pi_k = \frac{n_k}{N}$$

$$\mu_k = \frac{1}{n_k} \sum_{i=1}^N \gamma_i^k Y_i$$

$$\Sigma_k = \frac{1}{n_k} \sum_{i=1}^N \gamma_i^k (Y_i - \mu_k)(Y_i - \mu_k)'$$

# Implementing EM

```
gmm_loglik <- function(Y, K, mu, sigma, pi){  
  # Log-likelihood of each Y  
  lik_normal <- matrix(nrow=nrow(Y), ncol=K)  
  for(k in 1:K){  
    lik_normal[,k] <- mvtnorm::dmvnorm(Y, mu[k,], sigma[[k]])  
  }  
  # Log of the sums  
  log_likelihood <- sum(apply(lik_normal, 1, function(x) log(sum(x*pi))))  
  return(log_likelihood)  
}
```

# Implementing EM

```
gmm_estep_gamma <- function(Y, K, mu, sigma, pi){  
  # Calculate unnormalized gamma_k for each k  
  gamma <- matrix(nrow=nrow(Y), ncol=K)  
  for (k in 1:K){  
    gamma[,k] <- pi[k]*mvtnorm::dmvnorm(Y, mu[k,], sigma[[k]])  
  }  
  # Normalize the gammas (denominator)  
  gamma <- gamma/rowSums(gamma)  
  return(gamma)  
}
```



# Implementing EM

```
sigma_update <- function(Y, K, mu_k, gamma, n_k){ # Super inefficient but works
  # For each k
  sigma <- list()
  for (k in 1:K){
    sigma[[k]] <- matrix(data=0, nrow=ncol(Y), ncol=ncol(Y)) # blank matrix
    for (n in 1:nrow(Y)){
      sigma[[k]] <- sigma[[k]] + gamma[n,k]*outer(Y[n,] - mu_k[k,], Y[n,] - mu_k[k,])
    }
    sigma[[k]] <- sigma[[k]]/n_k[k]
  }
  return(sigma)
}
```

# Implementing EM

```
gmm_em <- function(Y, K, maxit=5000, tol=1e-8, verbose=F){
  D <- ncol(Y) # Number of dimensions
  ## Initialize the parameters (pick some reasonable starting points)
  mu_k <- MASS::mvrnorm(K, colMeans(Y), Sigma = var(Y)) # Matrix of means
  sigma_k <- list() # list of covariances
  for (k in 1:K){
    sigma_k[[k]] <- var(Y)
  }
  pi <- rep(1/K, K)
  curr_lik <- gmm_loglik(Y, K, mu_k, sigma_k, pi) #Evaluate current likelihood
  if(verbose) print(str_c("Log-Likelihood: ", curr_lik))
  for(iter in 2:maxit){
    # E-step
    gamma <- gmm_estep_gamma(Y, K, mu_k, sigma_k, pi)
    # M-step
    n_k <- colSums(gamma)
    pi <- n_k/sum(n_k)
    for(k in 1:K){
      mu_k[k,] <- colSums(gamma[,k]*Y)/n_k[k]
    }
    sigma_k <- sigma_update(Y, K, mu_k, gamma, n_k)
    # Check convergence
    new_lik <- gmm_loglik(Y, K, mu_k, sigma_k, pi)
    if (abs(new_lik - curr_lik) < tol){
      gamma <- gmm_estep_gamma(Y, K, mu_k, sigma_k, pi)
      if(verbose)print(str_c("Log-Likelihood: ", curr_lik))
      if(verbose) print(str_c("Log-Likelihood: ", new_lik))
    }
  }
}
```

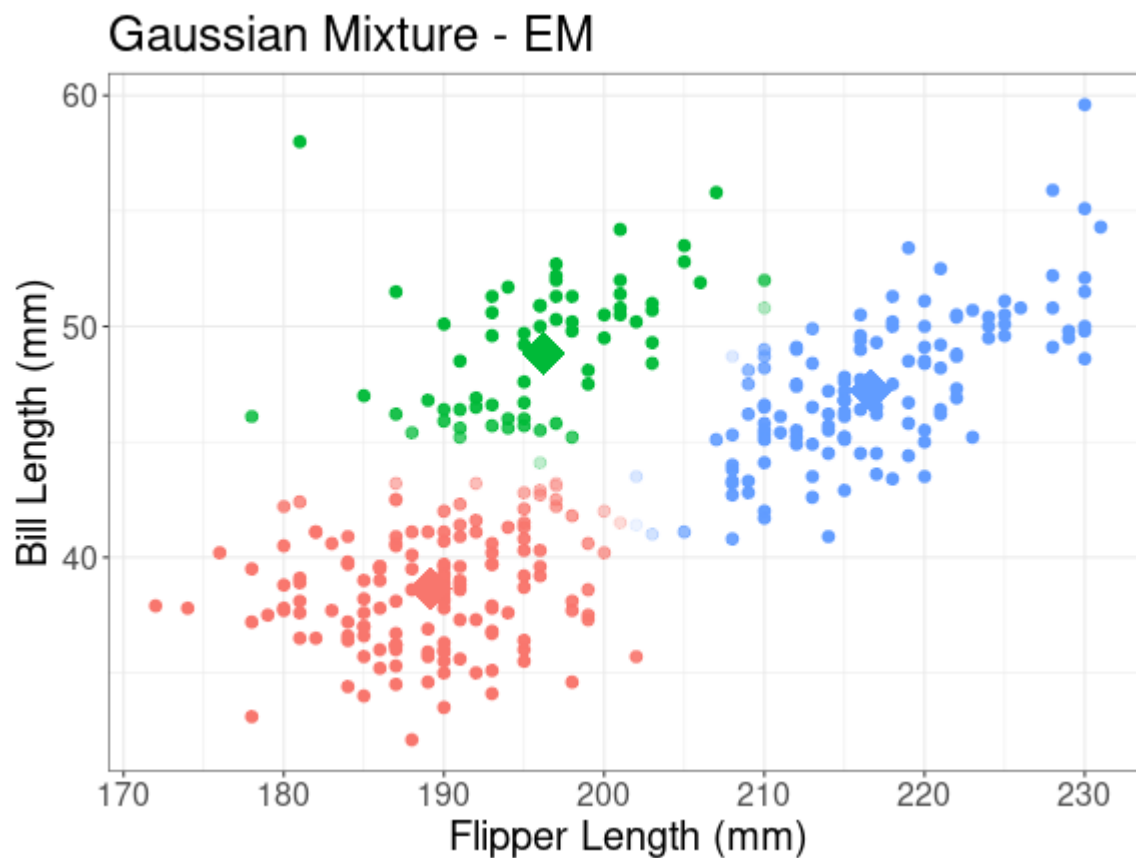
# Implementing EM

```
set.seed(60639)
penguins_em_3 <- gmm_em(Y=as.matrix(penguins_complete %>% select(flipper_length_mm, bill_length
```

```
## [1] "Log-Likelihood: -2464.23892416674"
## [1] "Log-Likelihood: -2286.49407632345"
## [1] "Log-Likelihood: -2265.16441563185"
## [1] "Log-Likelihood: -2244.45549391833"
## [1] "Log-Likelihood: -2244.22093530689"
## [1] "Log-Likelihood: -2244.21933979789"
## [1] "Log-Likelihood: -2244.21927829982"
## [1] "Log-Likelihood: -2244.21927596998"
## [1] "Log-Likelihood: -2244.21927591305"
## [1] "Log-Likelihood: -2244.21927590356"
## [1] "Complete!"
```

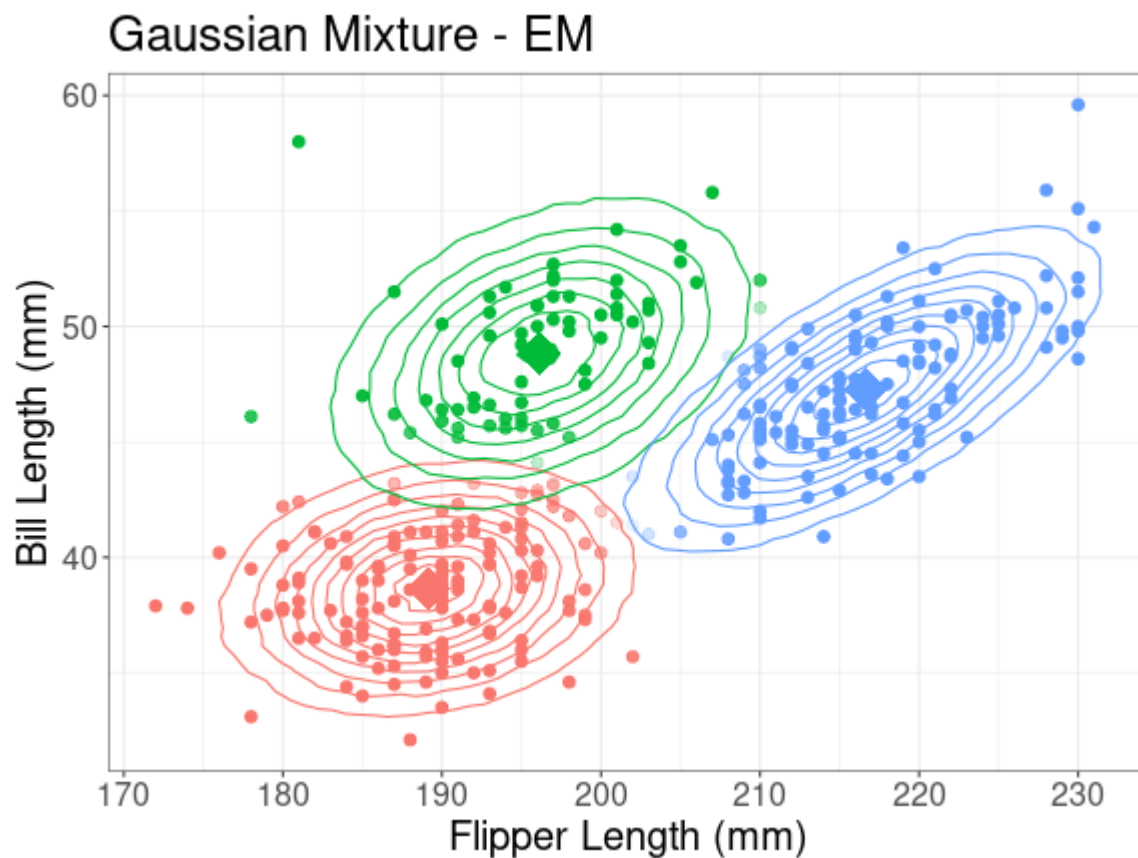
# Visualizing EM

- Now our point labels have **probabilities** attached




# Visualizing EM

- We can also visualize the component distributions



# Alternate K

- What does the 6-component mixture look like?



# Challenges with mixture models

- **Highly model dependent** - our identification of the cluster centroids + variances depends heavily on our selection of the appropriate distribution for  $Y_i$ 
  - Normal does okay for outcomes where it's plausible that CLT kicks in, but not true of all  $Y_i$
- **Multi-modality** - Mixture model likelihoods often have multiple modes and EM is only guaranteed to converge to a **local** optimum
  - Can get stuck in a "bad" EM run (**Solution**: Run multiple EM chains with different starting values and pick the one with the best log-likelihood)
  - "Label-switching" problem: Permuting the labels doesn't change the likelihood
- **Challenges with Bayes** - Sampling-based inference can be tricky with GMMs due to the label-switching problem.
  - Can implement many models via MCMC/Gibbs but need tricks to avoid having the chain jump between permutations.
  - Common to use an approximation to the likelihood around the posterior mode obtained via EM.
  - Stan doesn't like sampling discrete latent variables.


# Choosing $K$

- Choosing  $K$  is a problem of model comparison and selection.
  - Just as in regression, more parameters = better **in-sample** fit.
  - But we want to avoid **over-fitting**- Just as before, two approaches
  - **Information criteria** - Evaluate the in-sample log-likelihood at the maximum, penalized by some factor for the number of parameters
  - **Cross-validation** - Estimate the model on a **training set**. Compute the log-likelihood on a **held-out** test set.




# Choosing K

- Log-likelihood doesn't decrease as K increases, but common to look for an "elbow" where the increases become marginal




# Choosing K

- Information criteria results depend on how much you penalize the additional parameters
  - $AIC = 2k - 2\ell(\hat{\theta})$  shows minimal improvement beyond 6 clusters (but still potentially some!)



# Choosing K

- 10-fold cross-validation gives best performance at  $K = 3$



# Topic Modeling

# Modeling Text

- **Topic models** are a tool for **unsupervised** analysis of text
  - Don't need to label the texts ex-ante.
- Texts are **high-dimensional** - Topic models provide a **lower-dimensional** summary of the common themes in a document
  - Model document content as a **mixture** of a finite number of  $K$  latent **topics**
  - Topics are described as a **distribution** on words.
- History of topic models
  - **Latent Dirichlet Allocation**  $\rightsquigarrow$  **Correlated Topic Model**  $\rightsquigarrow$  **Structural Topic Model**

# Example: Political Blogs

- Our running example is the CMU 2008 Political Blog corpus (Eisenstein and Xing, 2010)
  - 13,246 posts from 8 political blogs rated liberal/conservative

```
library(stm)
poliblog <- read_csv("data/poliblogs2008.csv")
poliblog$documents[2:10] # Example document
```

```
## [1] "I honestly don't know how either party's caucus results will play out tonight. Usually, you can
## [2] "While we stand in awe of the willingness of our troops in Iraq to sacrifice themselves for the M
## [3] "These pages recently said goodbye to global warming.&nbsp;&nbsp;&nbsp;Ironically, the current spell of
## [4] "A US report shows how the enemy controlled the information on the battlefield in Fallujah and us
## [5] "Mike Huckabee is pretty slick. He's the one Republican who's been boosted by the big media. Thos
## [6] "In the aftermath of the \"upset\" victories in yesterday's Iowa caucus by Republican Mike Huckab
## [7] "FrontPage.Com has an editorial up this morning pointing out the deficiencies of both of the Iowa
## [8] "An influential policy outfit, the International Crisis Group, has called on President Musharraf
## [9] "Goodbye Iowa for 4 more years. But the Hawkeye State may have crowned the new President this tim
```

# Pre-processing the text

- We will need to convert the raw documents into a **document-term matrix**
  - **Rows**: Documents
  - **Columns**: "Tokens" (words)
- The `stm` package in R, which implements the **Structural Topic Model** has a host of helper functions for parsing texts and applying conventional pre-processing techniques (stemming, stop word removal, etc...)
  - `textProcessor` wraps the popular `tm` package for text mining

```
processed <- textProcessor(poliblog$documents, metadata=poliblog,  
                           lowercase=T, removestopwords=T,  
                           removenumbers = T, removepunctuation=T,  
                           stem = T)
```

```
out <- prepDocuments(processed$documents, processed$vocab, processed$meta,  
                      lower.thresh=round(.01*nrow(poliblog)), upper.thresh = round(.99*nrow(poliblog)))
```

```
## Removing 121334 of 123990 terms (531226 of 2298953 tokens) due to frequency  
## Your corpus now has 13246 documents, 2656 terms and 1767727 tokens.
```

# "Bag of Words"

- We summarize each document using the counts of each token
  - These constitute our **document-term matrix** (which software packages typically store in a "sparse" format)

```
head(data.frame(token=out$vocab[out$documents$`2`[1,]], count=out$documents$`2`[2,]))
```

##	token	count
## 1	ago	1
## 2	ahead	1
## 3	along	1
## 4	although	1
## 5	among	1
## 6	anoth	2


- The ordering of words is ignored - the "**bag of words**" assumption
  - Can preserve some ordering by tokenizing bigrams or trigrams rather than just unigrams (e.g. "republican party" or "daily tracking poll")
  - Interestingly, throwing out word order still typically preserves the underlying **topics** or **concepts** in a document!



# Latent Dirichlet Allocation

- **Latent Dirichlet Allocation** (Blei, Ng, Jordan, 2003)
  - A generative model for **words**  $w$  - defined as unit-vectors from a discrete vocabulary  $\{1, 2, \dots, V\}$
  - **Documents** are collections of words  $\mathbf{w} = \{w_1, w_2, \dots, w_N\}$
  - A **corpus** is a collection of documents  $\{\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_D\}$
- *LDA* assumes a particular generative process for each document  $\mathbf{w}_d$ :
  1. Choose a **topic proportion** vector  $\theta_d \sim \text{Dirichlet}(\alpha)$
  2. For each of the  $N$  words in document  $d$   $w_{d,n}$ 
    - Choose topic  $z_{d,n} \sim \text{Multinomial}(\theta_d)$
    - Choose word  $w_{d,n} \sim \text{Multinomial}(\beta_{z_{d,n}})$
- Two main targets of inference:
  - The **document topic proportions**:  $\theta = \{\theta_1, \theta_2, \dots, \theta_D\}$
  - The **topic-word distributions**:  $\beta = \{\beta_1, \beta_2, \dots, \beta_k\}$


# Latent Dirichlet Allocation



# Correlated Topic Model

- LDA makes a somewhat restrictive assumption on the distribution from where the topics are drawn (the Dirichlet)
  - Topic proportions are assumed to be **independent** (a feature of the dirichlet)
  - But in reality, some topics are more likely to appear with others!
- The **Correlated Topic Model** (Blei and Raftery, 2006) addresses this
  - Replaces the dirichlet with a logistic normal distribution.
- The **Correlated Topic Model** assumes that each document  $\mathbf{w}_d$  is generated by:
  1. Draw  $\eta_d$  from a **normal distribution** with mean  $\mu$ ,  $\Sigma$
  2. For each of the  $N$  words  $w_{d,n}$ 
    - Choose topic  $z_{d,n} \sim \text{Multinomial}(f(\eta_d))$  where  $f(\eta_{d,k}) = \exp(\eta_{d,k}) / \sum_j \exp(\eta_{d,j})$
    - Choose word  $w_{d,n} \sim \text{Multinomial}(\beta_{z_{d,n}})$


# Correlated Topic Model



# Structural Topic Model

- Often researchers want to incorporate **covariates** in modeling the topic proportions
  - e.g. - Do Liberal blogs talk about different topics compared to Conservative blogs?
  - Ad-hoc regressions of topic proportions on covariates will fail to incorporate uncertainty in the estimation of the topics.
- Additionally, researchers may want to have different word distributions for a topic depending on covariates
  - e.g. - Do Liberal blogs talk *about immigration* differently than Conservative blogs?
- The **Structural Topic Model** (Roberts, Stewart, Tingley, Airolidi, 2013) extends the correlated topic model to incorporate covariates:
  - Allow covariates  $X_d$  to enter into the logistic-normal distribution governing the topic proportions  $\eta_d \sim \text{Normal}(\mu_d, \Sigma)$  where  $\mu_{d,k} = X_d' \gamma_k$
  - Allow covariates  $Y_d$  to enter into the  $\beta$  parameter governing topic content
$$\beta_{d,k} \propto \exp(m + \kappa_k^{(t)} + \kappa_{y_d}^{(c)} + \kappa_{y_d,k}^{(i)})$$

# Structural Topic Model



# Illustration: Political Texts

- Without any covariates, `stm()` is a fast implementation of the **correlated topic model**

```
poliBlog_fit <- stm(documents = out$documents,  
                    vocab = out$vocab,  
                    K = 20, verbose=F,  
                    max.em.its = 75, data = out$meta,  
                    init.type = "Spectral", seed=60639)
```

- `stm()` implements a **Variational EM** algorithm to provide a tractable, analytical approximation to the posterior distribution
  - **Benefits:** Faster, deterministic
  - **Drawbacks:** Can understate the variance in the posterior
  - For more, see: Grimmer, 2010 "An Introduction to Bayesian Inference via Variational Approximations" *Political Analysis*

# Illustration: Political Texts

- Typically, we're interested in understanding the content of the topics discovered
  - We'll inspect the **topic-word** distributions to attempt to infer the underlying topic meaning from a set of "top" words

```
labelTopics(poliBlog_fit, topics=20)
```

```
## Topic 20 Top Words:  
## Highest Prob: financi, crisi, market, govern, money, bank, bailout  
## FREX: financi, bailout, loan, mortgag, bank, fanni, market  
## Lift: mae, fanni, paulson, freddi, subprim, treasuri, bailout  
## Score: mae, bailout, mortgag, fanni, financi, loan, market
```

- Common criteria for ranking the "top" words
  - **Highest probability** - Which words have the highest probability in the topic-word distribution  
 $\beta_k = p(w | z = k)$
  - **Frequency-Exclusivity** (FREX) - Harmonic mean of rank by within-topic probability and rank by topic distribution given word  $p(z | w = v)$
  - **Lift** - Topic-word distribution divided by the empirical distribution - how **more common** are these words in this topic relative to their baseline prevalence)
  - **Score** - Similar to Lift (but in terms of log probabilities)



# Illustration: Political Texts

- We can find "prototypical" documents that contain a high proportion of a particular topic

```
plotQuote(unlist(findThoughts(poliBlog_fit, texts = out$meta$documents, n=3, topics=20)$docs),
```

It would be the mother of all bailouts if it came to pass - hundreds of billions of taxpayer dollars used to shore up the two largest mortgage companies in the United States - Fannie Mae and Freddie Mac. The Federal National Mortgage Association (Fannie Mae) and the Federal Home Loan Mortgage

Over the weekend, the Federal Reserve opened a \$30 billion line of credit for the purchase of troubled investment bank Bear Sterns and promised an open ended lending program for the biggest investment firms on Wall Street: In a third move aimed at helping banks and thrifts, the Fed also lowered the

When the New York Times refers to the bailout plan for Citigroup as "radical," it must be somewhere out near Mars: Federal regulators approved a radical plan to stabilize Citigroup in an arrangement in which the government could soak up billions of dollars in losses at the struggling bank, the

# Illustration: Political Texts

- What about this topic?

```
labelTopics(poliBlog_fit, topics=16)
```

```
## Topic 16 Top Words:
```

```
## Highest Prob: will, american, care, health, america, peopl, can
```

```
## FREX: care, health, job, america, famili, educ, that
```

```
## Lift: coal, health, prosper, poverti, care, insur, afford
```

```
## Score: coal, health, tax, care, economi, american, insur
```

# Illustration: Political Texts

- Not all topics make sense

```
labelTopics(poliBlog_fit, topics=4)
```

```
## Topic 4 Top Words:
```

```
## Highest Prob: get, one, 're, don't, like, want, doesn't
```

```
## FREX: doesn't, 'll, 're, don't, didn't, can't, 've
```

```
## Lift: chat, see-dubya, 'll, one', can't, isn't, doesn't
```

```
## Score: chat, 're, 'll, don't, doesn't, 've, didn't
```

# Illustration: Political Texts

- Topic 4 seems to be picking up some odd stylistic feature of Minnesota conservative talk radio?

```
plotQuote(unlist(findThoughts(poliBlog_fit, texts = out$meta$documents, n=3, topics=4)$docs), n
```

The Northern Alliance Radio Network will be on the air today, with our eight-hour-long broadcast schedule starting at 9 am CT. If you're in the Twin Cities, you can hear us on AM 1280 The Patriot, or on the station's Internet stream if you're outside of the broadcast area. Today, I'm traveling to

---

I woke up this morning to a flood of e-mails, some angry and some apologetic, from people wondering why they'd been summarily booted from the site without so much as a word of explanation. The answer: You weren't. Our web wizard, Mark Jaquith, upgraded HA to the newest version of WordPress last

---

The Northern Alliance Radio Network will be on the air today, with our eight-hour-long broadcast schedule starting at 9 am CT. If you're in the Twin Cities, you can hear us on AM 1280 The Patriot, or on the station's Internet stream if you're outside of the broadcast area. Today, Mitch and I will

# Choosing Models

- How do we evaluate the quality of a model?
  - ...for multiple runs of the same topic count (multiple modes)
  - ...for different numbers of topics?
- Conventional method: **held-out log-likelihood**. For an "out-of-sample" document  $\mathbf{w}_{\text{out}}$ :

$$\text{Perplexity} = \exp(-\log p(\mathbf{w}_{\text{out}} | \mu, \Sigma, \beta))$$

- But "highly predictive" topic models may be difficult for humans to interpret (Chang et. al., 2009).

# Coherence and Exclusivity

- **Alternative** - Pick models where the topics have nice features
  - **Semantic coherence** - The most common words in a topic tend to co-occur in a document
  - **Exclusivity** - Words that have high probability in a topic have low probability in **other** topics (FREX).
- **Semantic coherence**
  - Let  $D(v, v')$  define the number of times words  $v$  and  $v'$  appear together in a document.
  - Semantic coherence is defined as a sum across the  $W$  most probable words

$$C_k = \sum_{i=2}^W \sum_{j=1}^{i-1} \log \left( \frac{D(v_i, v_j) + 1}{D(v_j)} \right)$$

# Illustration: Political Texts

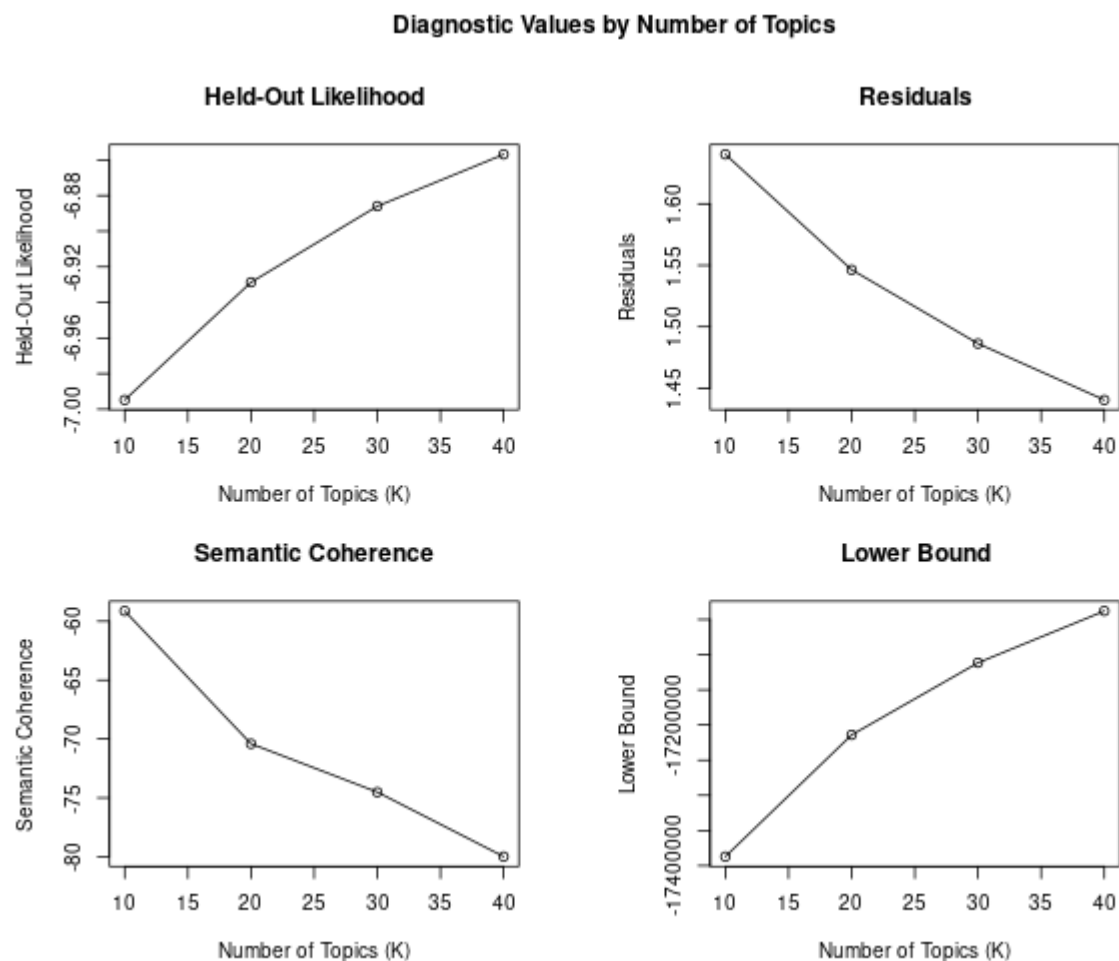
- Let's compare fits across 10, 20, 30, and 40 topics

```
modelEval <- searchK(out$documents,  
                     vocab = out$vocab,  
                     K = c(10,20,30,40), verbose=F,  
                     max.em.its = 75, data = out$meta,  
                     init.type = "Spectral", seed=60639, cores=4)
```

## Using multiple-cores. Progress will not be shown.


# Illustration: Political Texts

- Average topic exclusivity





# Illustration: Political Texts



# Estimating coefficients

- With `stm()`, we can incorporate covariates into the topic proportions using standard formula syntax

```
poliBlog_fit_cov <- stm(documents = out$documents,  
                        vocab = out$vocab,  
                        prevalence = ~ rating,  
                        K = 20, verbose=F,  
                        max.em.its = 75, data = out$meta,  
                        init.type = "Spectral", seed=60639)
```

# Estimating coefficients

- Do liberal blogs talk about health care more than conservative blogs?

```
estim_hc <- estimateEffect(c(16) ~ rating, stmobj=poliBlog_fit_cov, metadata=out$meta)
summary(estim_hc)
```

```
##
## Call:
## estimateEffect(formula = c(16) ~ rating, stmobj = poliBlog_fit_cov,
##               metadata = out$meta)
##
##
## Topic 16:
##
## Coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.03106    0.00117   26.6   <2e-16 ***
## ratingLiberal  0.02166    0.00182   11.9   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

# Estimating coefficients

- Do conservative blogs talk about the financial crisis more than liberal blogs?

```
estim_fc <- estimateEffect(c(20) ~ rating, stmobj=poliBlog_fit_cov, metadata=out$meta)
summary(estim_fc)
```

```
##
## Call:
## estimateEffect(formula = c(20) ~ rating, stmobj = poliBlog_fit_cov,
##   metadata = out$meta)
##
##
## Topic 20:
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)   0.04007    0.00130   30.79  <2e-16 ***
## ratingLiberal -0.00547    0.00210   -2.61   0.0091 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

